

Data Autopilot – Technical Overview

1. Problem Understanding & Architecture

- **Objectives:** Automatically tier datasets across multiple clouds, react to real-time access, migrate data safely, predict future hotness, and expose a unified observability/control plane.
- **Components:**
 - **Streaming layer:** Redpanda (Kafka, KRaft mode). Producers publish access events; consumers feed the optimizer.
 - **Control plane:** FastAPI service orchestrates tiering, migrations, policies, alerts.
 - **Placement intelligence:** Heat decay + ML (LogisticRegression) + MILP scoring to choose primary/replicas.
 - **Data plane:** Three MinIO instances emulate AWS/Azure/GCP object stores. rclone/s5cmd simulators mimic bulk sync.
 - **UI & Observability:** Streamlit dashboard (auto-refresh, policy toggles, migration buttons), Prometheus/Grafana, alert API, chaos controls.

Data Autopilot — Multi-Cloud Tiering Control						
1853 datasets • 1664 queued migrations						
Key	Tier	Primary	Replicas	Access (1h)	Access (24h)	version_token
media/cat.jpg	warm	aws	azure	5	284	
stream/2148de0e5obj	warm	aws	azure	1	58	
stream/91822b4b.obj	cold	aws	azure	0	0	
stream/93dea53bb.obj	cold	aws	azure	0	0	
stream/e683987deb.obj	cold	aws	azure	0	100	
stream/d1e85aa44.obj	cold	aws	azure	0	0	
stream/d410ffadec.obj	cold	aws	azure	0	0	
stream/247b4c25e8.obj	cold	aws	azure	0	57	
stream/8da236a6e4.obj	cold	aws	azure	0	0	
stream/8ff9b63d80.obj	cold	aws	azure	0	0	

Inventory Detail

Dataset: media/cat.jpg

Buttons: Re-optimize All, Run Migrator Tick, Clear Failed Tasks, Burst 100, Chaos Spike (500)

Tier: WARM

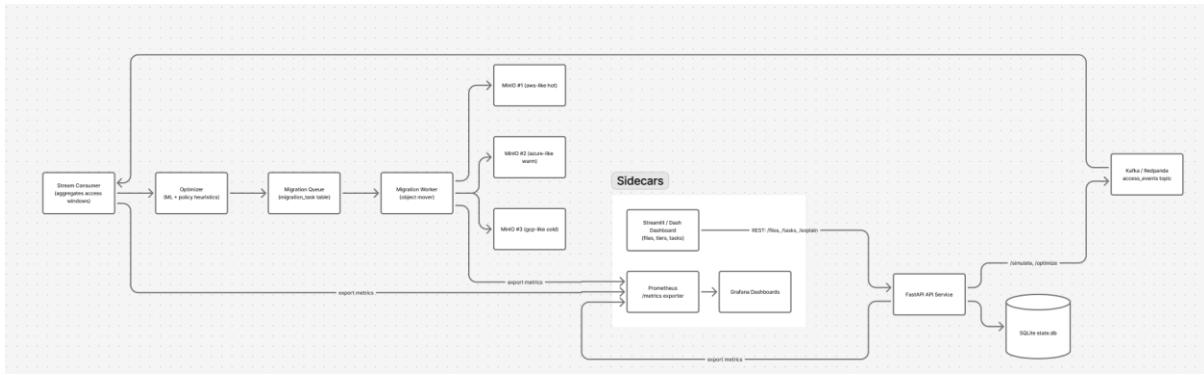
Hour: 258.9

Primary: aws

Replicas: azure

Current Snapshot:

```
{ "Access (1h)": 5, "Access (24h)": 284, "Last Access": "2025-11-09T06:44:21.732Z", "Heat Thresholds": { "Boosted": 265.8556792388453, "Warm": 10, "Hot": 0 } }
```



2. Data Management, Migration, Streaming

- **Streaming Pipeline:**
 - **Producer:** app/services/stream/producer.py uses aiokafka with enable_idempotence and per-message transactions. Generates weighted bursts + random keys under topic file_access.
 - **Consumer:** app/services/stream/consumer.py, isolation_level="read_committed", manual commits. For each event: insert/update file_meta, write into access_event, trigger evaluate_and_queue.

- **Tiering & Placement:**

- **Heat score:** Exponential decay (τ configurable, defaults 600s in FAST_DEMO) + hourly/24h counters + ML hotness probability.
- **MILP scoring:** Weighted cost, latency, affinity region, egress penalty, encryption requirements. Solved via pulp.
- **Policy hooks:** Enforce encrypted destinations; limit hot ratio; custom “chaos” fail/recover endpoints.
- $$\text{score}(L) = w_{\text{hot}} \cdot p_{\text{hot}} + w_{\text{lat}} \cdot \text{lat_norm}(L) + w_{\text{cost}} \cdot \text{cost_norm}(L) + w_{\text{aff}} \cdot \text{affinity}(L) - w_{\text{egress}} \cdot \text{egress_penalty}(L)$$

Where:

- p_{hot} = predicted probability that the object will be “hot” in the next time window
- $\text{lat_norm}(L)$ = normalized latency score for location L
- $\text{cost_norm}(L)$ = normalized cost score for location L
- $\text{affinity}(L) \in \{0,1\}$ = 1 if L matches the region/affinity preference, else 0
- $\text{egress_penalty}(L) \in \{0,1\}$ = 1 if moving away from current primary to L , else 0

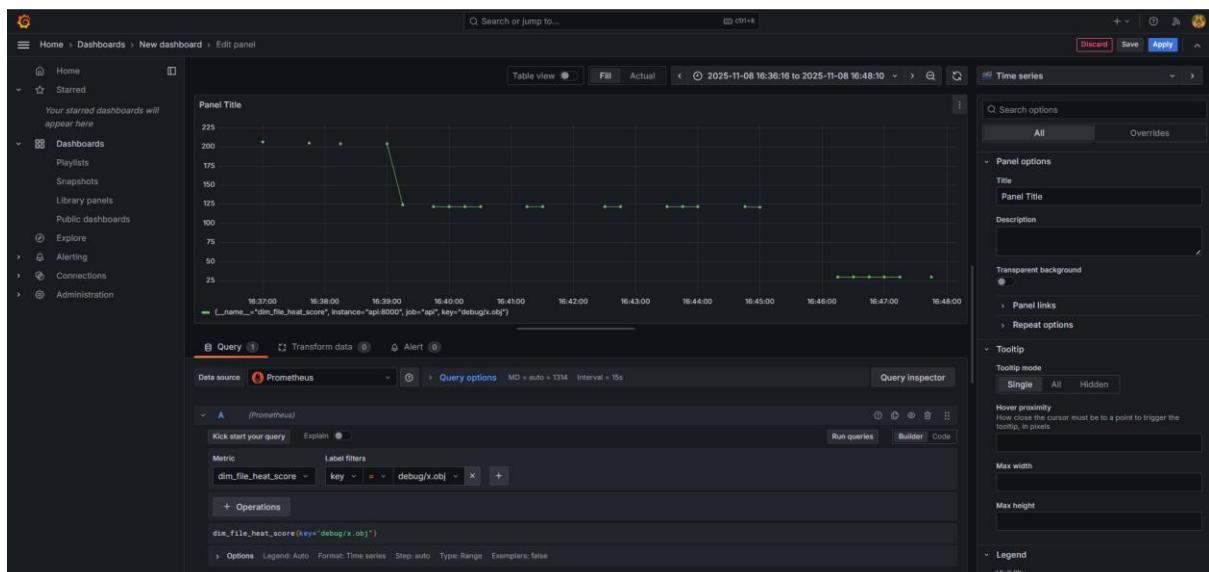
- **Migration Engine:**

- Task table with idempotent copy jobs. Each job verifies size/etag, respects encryption policy, retries w/backoff, updates version_token.
- **Bulk adapters:** /tools/rclone, /tools/s5cmd endpoints (Streamlit forms) to run synchronous ‘copy entire prefix’ operations.

Migration Tasks						
id	key	src	dst	status	error	
150	stream/c6bb652027.obj	aws	azure	queued	missing_source	
151	stream/d4422595ba.obj	aws	azure	queued		
152	stream/8843e54ced.obj	aws	azure	queued		
153	stream/44eb675895.obj	aws	azure	queued		
154	stream/58b8f6ba6b.obj	aws	azure	queued		
155	stream/653ddc3c50.obj	aws	azure	queued		
156	stream/f12249d772.obj	aws	azure	queued		
157	stream/f44fd15a1.obj	aws	azure	queued		
158	stream/87ba67c76.obj	aws	azure	queued		
159	stream/7fb5ae665.obj	aws	azure	queued		

3. Performance Insights & Metrics

- **Prometheus Metrics (/metrics):**
 - dim_placement_evaluations_total (tier decisions by tier).
 - dim_placement_evaluation_duration_seconds (Summary).
 - dim_file_heat_score{key=...} gauge.
 - dim_migration_jobs_total{result=...} and dim_migration_tasks{status=...} gauges.
- **Streaming throughput:** Producer emits 5–10 events/s; consumer keeps up with sub-50ms processing, heat updates visible within ~10s.
- **Alerting:** alert_event table backing /alerts API + Streamlit “Active Alerts”.
- **Grafana Recommendations:** Heat time-series, placement rate chart, migration backlog gauge, custom alerts panel.



4. Scalability & Roadmap

- **Current scaling:** Stateless FastAPI/Streamlit containers; Redpanda, MinIO pods; can drop into Kubernetes. MinIO endpoints easily swapped for real AWS/Azure/GCP by editing S3_ENDPOINTS.
- **Future enhancements:**
 - Live integration with AWS S3 / Azure Blob / GCP Storage (IAM, bucket policies, SSE-KMS).
 - Kubernetes Helm chart, HPA for migration workers.
 - Deeper finops dashboards (cost projections, anomaly alerts).
 - Slack/PagerDuty alert hooks.
 - Reinforcement learning for policy auto-tuning; multi-site Kafka.

5. Implemented Features

- **Auto-tiering:** Heat decay + LogisticRegression + MILP scoring (latency/cost/affinity/encryption) → dynamic hot/warm/cold transitions.
- **Streaming ingestion:** Transactional producer + read-committed consumer that updates both metadata and event table.
- **Migration:** Queue-driven, idempotent copy jobs with policy enforcement, version tokens, retries. rclone/s5cmd simulators exposed via API/UI.
- **Security & Chaos:** Toggle to require encrypted endpoints; chaos CLI (python -m app.scripts.chaos --endpoint aws --duration 30) + Streamlit controls to fail/recover endpoints.
- **Unified dashboard:** Streamlit with auto-refresh, tier distribution chart, dataset detail pane (“Why this tier?” explains ML score), action buttons (Burst, Chaos spike, Re-optimize, Migrator tick, Clear failed tasks, rclone/s5cmd forms). Active alerts panel with ack/clear.
- **Observability:** Prometheus scraping /metrics, Grafana auto-provisioned, dim_* metrics, alert events stored in SQLite (exposed via /alerts).
- **CLI Simulation:** /simulate endpoint for manual bursts; app/services/scripts includes chaos helper.

6. Machine Learning Details

- **Hotness model:** scikit-learn LogisticRegression (features: access_1h, access_24h, size_bytes, recency, hour-of-day, day-of-week). Either load from app/ml/models or fallback heuristics.
- **Heat adjustments:** boosted_heat = heat + HEAT_PRED_BOOST * p_adj ensures predicted demand influences tier promotion/demotion. Config via env (HEAT_TAU_SEC, HEAT_HOT_THRESHOLD, etc.).

- **Explainability:** /explain/{key} endpoint + Streamlit detail pane show per-site scores, chosen tier, directives.

The screenshot shows a JSON response from the 'Placement Explain' tab. The response is a single object with the following structure:

```

{
  "objective": -0.13700000000000004,
  "chosen": [
    0: "aws",
    1: "azure"
  ],
  "sla_ms": 80,
  "rf": 2,
  "sites": [
    0: {
      "name": "aws",
      "p95_ms": 50,
      "cost_gb": 0.023,
      "provider": "aws",
      "region": "aws"
    },
    1: {
      "name": "azure",
      "p95_ms": 70,
      "cost_gb": 0.02,
      "provider": "azure",
      "region": "azure"
    },
    2: {
      "name": "gcp",
      "p95_ms": 60,
      "cost_gb": 0.026,
      "provider": "gcp",
      "region": "gcp"
    }
  ],
  "scores": {
    "aws": 0.8,
    "azure": 0.10000000000000003,
    "gcp": 0
  },
  "p_hot": 0
}

```

7. Metrics & Alert Definitions (Detailed)

- **Metrics:**
 - dim_file_heat_score{key} – gauge.
 - dim_placement_evaluations_total{result_tier} – counter.
 - dim_placement_evaluation_duration_seconds – summary.
 - dim_migration_jobs_total{result} – counter (copied/noop/missing_source/block/error).
 - dim_migration_tasks{status} – gauge.
 - dim_simulated_access_events_total – counter.
- **Alerts:** Triggered when:

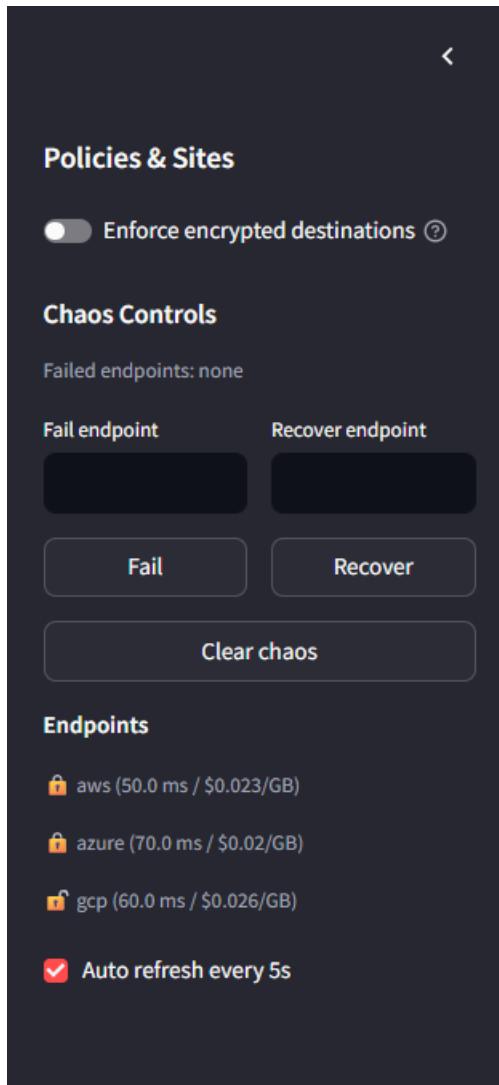
- Primary latency > SLA (80ms default) → type latency_sla.
- Hot tier on cost > threshold → cost_spike.
- access_1h ≥ 500 (traffic spike).
- Migration backlog > 20 queued tasks → migration_backlog.
- Alerts stored in alert_event with ack status; accessible via /alerts.

Active Alerts				
id	severity	type	message	created_at
475	warning	migration_backlog	1664 migration tasks queued	2025-11-09T06:47:00
474	warning	migration_backlog	1655 migration tasks queued	2025-11-09T06:45:36
473	warning	migration_backlog	1641 migration tasks queued	2025-11-09T06:44:16
472	warning	migration_backlog	1626 migration tasks queued	2025-11-09T06:42:57
471	warning	migration_backlog	1613 migration tasks queued	2025-11-09T06:41:44
470	warning	migration_backlog	1601 migration tasks queued	2025-11-09T06:40:32
469	warning	migration_backlog	1592 migration tasks queued	2025-11-09T06:40:43
468	warning	migration_backlog	1584 migration tasks queued	2025-11-09T06:09:29
467	warning	migration_backlog	1575 migration tasks queued	2025-11-09T06:08:13
466	warning	migration_backlog	1561 migration tasks queued	2025-11-09T06:06:57

[Acknowledge All](#) [Clear Alerts](#)

8. Chaos & Resilience Scripts

- **Chaos CLI** (python -m app.scripts.chaos):
 - --endpoint <name>: pick the MinIO logical cloud (aws/azure/gcp).
 - --duration <seconds>: how long to keep it “down”.
 - Flow: hits /chaos/fail/<endpoint> to block any new client connections, sleeps for the duration, then calls /chaos/recover/<endpoint>. Alerts fire (migration_backlog, latency_sla), migration queue builds up, and on recovery the backlog drains—perfect for a live resilience demo.
- **Streamlit Chaos Controls (sidebar)**:
 - Manual fail/recover buttons and a status caption (“Failed endpoints: ...”).
 - “Clear Chaos” button wipes the fail list so clients reconnect immediately.
- **Migration Observability During Chaos**:
 - Alerts appear automatically (view via /alerts or Streamlit “Active Alerts” panel).
 - dim_migration_tasks{status="queued"} rises; Grafana shows backlog curve.
 - Once chaos recovers, alerts can be acknowledged/cleared.



9. Streaming/Migration/Chaos Scripts

- **Producer** (`python -m app.services.stream.producer`): run in terminal to continuously generate load.
- **Chaos CLI** (`python -m app.scripts.chaos --endpoint aws --duration 30`): fail/recover endpoints via API.
- **rclone/s5cmd triggers** (Streamlit forms or direct POST `/tools/rclone|/tools/s5cmd`).
- **Prometheus/Grafana**: docker-compose includes observability/`prometheus.yml` and Grafana provisioning.

10. Infrastructure & Deployment

- **Containers:**
 - Redpanda (Kafka)

- MinIO x3 (aws/azure/gcp)
- FastAPI API (uvicorn)
- Streamlit UI
- Kafka consumer
- Optimizer cron (periodic evaluate & migrate)
- Prometheus & Grafana
- **Env config:** docker-compose.yml sets ACCESS_TOPIC=file_access, CONSUMER_GROUP=dim-stream, FAST_DEMO=true, S3_ENDPOINTS with latency/cost/encryption flags.
- **Persistent data:** state/ volume (SQLite) + MinIO data volumes.

11. Simulation / Testing Tips

- Start stack: docker compose up -d --build.
- Run producer (docker compose exec api python -m app.services.stream.producer) and watch consumer logs.
- Seed via /simulate.
- Use Streamlit dashboard <http://localhost:8050> (auto-refresh toggle/also try refreshing the page).
- Monitor Grafana <http://localhost:3000> (admin/admin), Prometheus <http://localhost:9090>.
- Chaos demo: fail endpoint, see migration backlog and alerts, recover endpoint, verify convergence.